
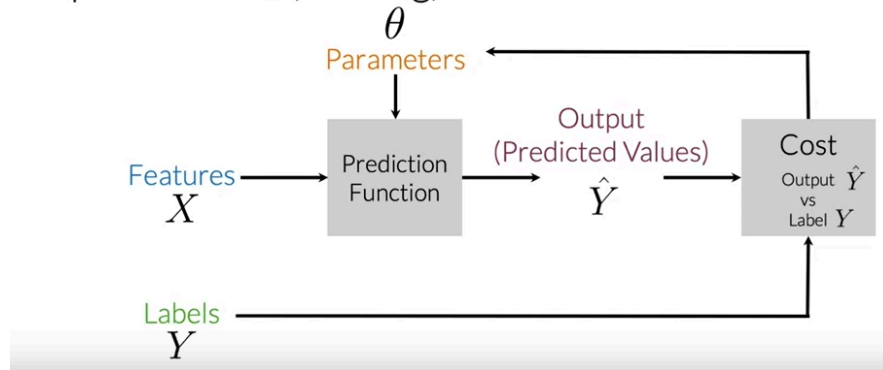


# Classification and Vector Spaces:

## Supervised ML:

Supervised ML (training)



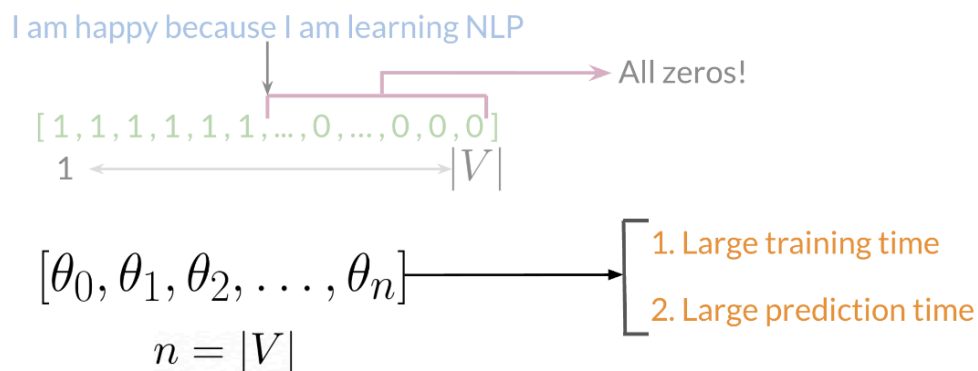
## Vocabulary (V):

Set of all possible words in the corpus

## Feature Extraction:

### Sparse Representation:

For each word, assign 1 if present in vocab, else assign 0



## Positive and Negative Frequencies:

For each word in the vocabulary, count how many times that word appear in the corpus of that particular class (positive or negative).

### Positive tweets

I am happy because I am learning NLP  
I am happy

### Negative tweets

I am sad, I am not learning NLP  
I am sad

Vocabulary	PosFreq (1)	NegFreq (0)
I	3	3
am	3	3
happy	2	0
because	1	0
learning	1	1
NLP	1	1
sad	0	2
not	0	1

### Feature Extraction with Frequencies:

I am sad, I am not learning NLP

I am sad, I am not learning NLP

$$X_m = [1, \sum_w \text{freqs}(w, 1), \sum_w \text{freqs}(w, 0)]$$

↓  
8

$$X_m = [1, \sum_w \text{freqs}(w, 1), \sum_w \text{freqs}(w, 0)]$$

↓  
11

Here summation is the sum of freq of "set of words" for that sentence.

### Preprocessing:

When preprocessing, you have to perform the following:

1. Eliminate handles and URLs
2. Tokenize the string into words.
3. Remove stop words like "and, is, a, on, etc."
4. Stemming- or convert every word to its stem. Like dancer, dancing, danced, becomes 'danc'. You can use porter stemmer to take care of this.
5. Convert all your words to lower case.

### 6. Removing Punctuations

Ex:

@YMcourri and @AndrewYNg are tuning a GREAT AI model at <https://deeplearning.ai!!!>

Preprocessed tweet:


[tun, great, ai, model]

### Complete Implementation:

```
freqs = build_freqs(tweets, labels) #Build frequencies dictionary
X = np.zeros((m, 3)) #Initialize matrix X
for i in range(m): #For every tweet
    p_tweet = process_tweet(tweets[i]) #Process tweet
    X[i, :] = extract_features(p_tweet, freqs) #Extract Features
```

### Logistic Regression:

$$h(x^{(i)}, \theta) = \frac{1}{1 + e^{-\theta^T x^{(i)}}}$$




It uses sigmoid function

@YMurri and  
@AndrewYNg are tuning a  
GREAT AI model

[tun, ai, great, model]


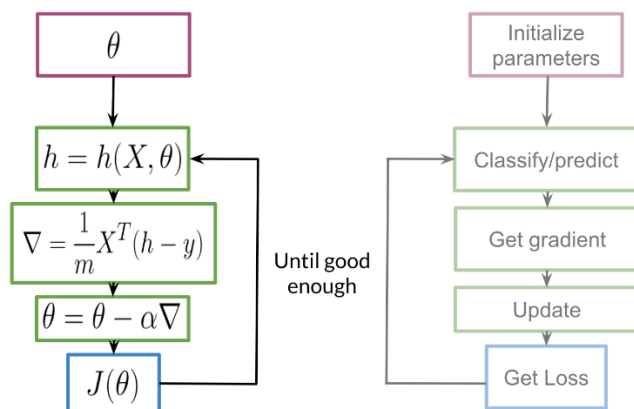
$$x^{(i)} = \begin{bmatrix} 1 \\ 3476 \\ 245 \end{bmatrix} \quad \theta = \begin{bmatrix} 0.00003 \\ 0.00150 \\ -0.00120 \end{bmatrix}$$



## Training Classifier:

### Logistic Regression: Training

To train your logistic regression function, you will do the following:



## Testing Classifier:

•  $X_{val} \ Y_{val} \ \theta$

$$h(X_{val}, \theta)$$

$$pred = h(X_{val}, \theta) \geq 0.5$$


$$\begin{bmatrix} 0.3 \\ 0.8 \\ 0.5 \\ \vdots \\ h_m \end{bmatrix} \geq 0.5 = \begin{bmatrix} 0.3 \geq 0.5 \\ 0.8 \geq 0.5 \\ 0.5 \geq 0.5 \\ \vdots \\ pred_m \geq 0.5 \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \\ 1 \\ \vdots \\ pred_m \end{bmatrix}$$

$$\text{Accuracy} \longrightarrow \sum_{i=1}^m \frac{(pred^{(i)} == y_{val}^{(i)})}{m}$$


### Cost Function:

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log (1 - h(x^{(i)}, \theta))]$$

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log (1 - h(x^{(i)}, \theta))]$$



**Y = 1**



**Y = 0**

### Derivation of loss function:

<https://www.coursera.org/learn/classification-vector-spaces-in-nlp/supplement/b3fHH/optimal-logistic-regression-cost-function>


### Derivation of gradient:

<https://www.coursera.org/learn/classification-vector-spaces-in-nlp/supplement/afcaR/optimal-logistic-regression-gradient>

### Naive Bayes:

### Probability:


### Corpus of tweets




$A \rightarrow$  Positive tweet

$$P(A) = N_{\text{pos}} / N = 13 / 20 = 0.65$$

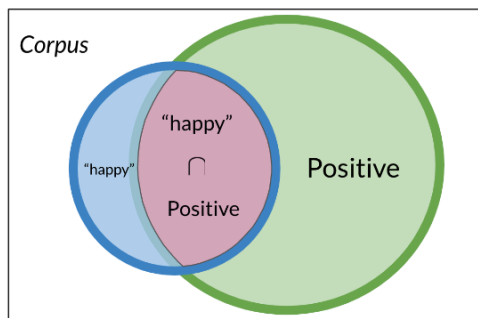
$$P(\text{Negative}) = 1 - P(\text{Positive}) = 0.35$$



$$P(A \cap B) = P(A, B) = \frac{3}{20} = 0.15$$



### Conditional Probabilities:



$$P(\text{Positive} | \text{"happy"}) =$$

$$\frac{P(\text{Positive} \cap \text{"happy"})}{P(\text{"happy"})}$$

### Bayes Rule:

$$P(\text{Positive} | \text{"happy"}) = \frac{P(\text{Positive} \cap \text{"happy"})}{P(\text{"happy"})}$$

$$P(\text{"happy"} | \text{Positive}) = \frac{P(\text{"happy"} \cap \text{Positive})}{P(\text{Positive})}$$

$$P(\text{Positive} | \text{"happy"}) = P(\text{"happy"} | \text{Positive}) \times \frac{P(\text{Positive})}{P(\text{"happy"})}$$

$$P(X|Y) = \frac{P(Y|X)P(X)}{P(Y)}$$

## Naive Bayes:

Positive tweets
I am happy because I am learning NLP
I am happy, not sad.
Negative tweets
I am sad, I am not learning NLP
I am sad, not happy

word	Pos	Neg
I	3	3
am	3	3
happy	2	1
because	1	0
learning	1	1
NLP	1	1
sad	1	2
not	1	2
<b>N<sub>class</sub></b>	<b>13</b>	<b>12</b>

word	Pos	Neg
I	0.24	0.25
am	0.24	0.25
happy	0.15	0.08
because	0.08	0
learning	0.08	0.08
NLP	0.08	0.08
sad	0.08	0.17
not	0.08	0.17

P(word|pos\_class) or P(word|neg\_class)

Tweet: I am happy today; I am learning.

$$\prod_{i=1}^m \frac{P(w_i|pos)}{P(w_i|neg)} = \frac{0.14}{0.10} = 1.4 > 1$$

$$\frac{\cancel{0.20}}{0.20} * \frac{\cancel{0.20}}{0.20} * \frac{0.14}{0.10} * \frac{\cancel{0.20}}{0.20} * \frac{\cancel{0.20}}{0.20} * \frac{\cancel{0.10}}{0.10}$$

word	Pos	Neg
I	0.20	0.20
am	0.20	0.20
happy	0.14	0.10
because	0.10	0.05
learning	0.10	0.10
NLP	0.10	0.10
sad	0.10	0.15
not	0.10	0.15

Bayes

Naive

## Laplacian Smoothing:

We usually compute the probability of a word given a class as follows:

$$P(w_i | \text{class}) = \frac{\text{freq}(w_i, \text{class})}{N_{\text{class}}} \quad \text{class} \in \{ \text{Positive}, \text{Negative} \}$$

However, if a word does not appear in the training, then it automatically gets a probability of 0, to fix this we add smoothing as follows

$$P(w_i | \text{class}) = \frac{\text{freq}(w_i, \text{class}) + 1}{(N_{\text{class}} + V)}$$

Note that we added a 1 in the numerator, and since there are  $V$  words to normalize, we add  $V$  in the denominator.

$N_{\text{class}}$ : frequency of all words in class

$V$ : number of unique words in vocabulary

## Log Likelihood:

$$\log \left( \frac{P(\text{pos})}{P(\text{neg})} \prod_{i=1}^n \frac{P(w_i|\text{pos})}{P(w_i|\text{neg})} \right) \Rightarrow \log \frac{P(\text{pos})}{P(\text{neg})} + \sum_{i=1}^n \log \frac{P(w_i|\text{pos})}{P(w_i|\text{neg})}$$

The first component is called the log prior and the second component is the log likelihood.

doc: I am happy because I am learning.

$$\lambda(w) = \log \frac{P(w|\text{pos})}{P(w|\text{neg})}$$

$$\lambda(\text{happy}) = \log \frac{0.09}{0.01} \approx 2.2$$

word	Pos	Neg	$\lambda$
I	0.05	0.05	0
am	0.04	0.04	0
happy	0.09	0.01	
because	0.01	0.01	
learning	0.03	0.01	
NLP	0.02	0.02	
sad	0.01	0.09	
not	0.02	0.03	

doc: I am happy because I am learning.

$$\sum_{i=1}^m \log \frac{P(w_i|\text{pos})}{P(w_i|\text{neg})} = \sum_{i=1}^m \lambda(w_i)$$

$$\log \text{likelihood} = 0 + 0 + 2.2 + 0 + 0 + 0 + 1.1 = 3.3$$

word	Pos	Neg	$\lambda$
I	0.05	0.05	0
am	0.04	0.04	0
happy	0.09	0.01	2.2
because	0.01	0.01	0
learning	0.03	0.01	1.1
NLP	0.02	0.02	0
sad	0.01	0.09	-2.2
not	0.02	0.03	-0.4

## Training Naive Bayes:

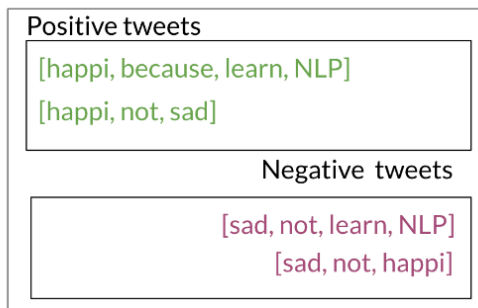


**1) Get or annotate a dataset with positive and negative tweets**

**2) Preprocess the tweets: `process_tweet(tweet) → [w1, w2, w3, ...]`:**

- Lowercase
- Remove punctuation, urls, names
- Remove stop words
- Stemming
- Tokenize sentences

**3) Compute `freq(w, class)`:**



word	Pos	Neg
happi	2	1
because	1	0
learn	1	1
NLP	1	1
sad	1	2
not	1	2
$N_{\text{class}}$	7	7

freq(w, class)

**4) Get  $P(w|pos)$ ,  $P(w|neg)$**

You can use the table above to compute the probabilities.

**5) Get  $\lambda(w)$**

$$\lambda(w) = \log \frac{P(w|pos)}{P(w|neg)}$$

**6) Compute  $\logprior = \log(P(pos)/P(neg))$**

$\logprior = \log \frac{D_{pos}}{D_{neg}}$ , where  $D_{pos}$  and  $D_{neg}$  correspond to the number of positive and negative documents respectively.

**Testing Naive Bayes:**

- log-likelihood dictionary  $\lambda(w) = \log \frac{P(w|pos)}{P(w|neg)}$

- $logprior = \log \frac{D_{pos}}{D_{neg}} = 0$

- Tweet: [I, pass, the NLP interview] 🍀

$$score = -0.01 + 0.5 - 0.01 + 0 + logprior = 0.48$$

$$pred = score > 0$$

word	$\lambda$
I	-0.01
the	-0.01
happi	0.63
because	0.01
pass	0.5
NLP	0
sad	-0.75
not	-0.75

### Application:

- Author identification
- Spam filtering
- Information retrieval
- Word disambiguation

### Assumptions:

- Independence
- Relative frequency in corpus

Means unequal frequencies of classes in the dataset

### Errors:

- Removing punctuation and stop words
- Word order
- Adversarial attacks

### Vector Space:

#### Word by Word Design:

I like simple data

I prefer simple raw data


k=2

	simple	raw	like	I
data	2	1	1	0

n

## Word by document design:

	Entertainment	Economy	Machine Learning
data	500	6620	9320
film	7000	4000	1000



	Entertainment	Economy	ML
data	500	6620	9320
film	7000	4000	1000


Measures of "similarity:"  
Angle  
Distance

## Euclidean Distance:

$$d(B, A) = \sqrt{((B_1 - A_1)^2 + (B_2 - A_2)^2)}$$

$$d(\vec{v}, \vec{w}) = \sqrt{\sum_{i=1}^n (v_i - w_i)^2}$$

## Cosine Similarity:




Euclidean distance:  $d_2 < d_1$

Angles comparison:  $\beta > \alpha$

The cosine of the angle  
between the vectors

If corpus are of different sizes, then cosine similarity is a better metric.




$$\hat{v} \cdot \hat{w} = \|\hat{v}\| \|\hat{w}\| \cos(\beta)$$

$$\cos(\beta) = \frac{\hat{v} \cdot \hat{w}}{\|\hat{v}\| \|\hat{w}\|}$$

$$= \frac{(20 \times 30) + (40 \times 20)}{\sqrt{20^2 + 40^2} \times \sqrt{30^2 + 20^2}} = 0.87$$

### Manipulating vectors:



$$\text{Washington} - \text{USA} = \begin{bmatrix} 5 & -1 \end{bmatrix}$$

$$\text{Russia} + \begin{bmatrix} 5 & -1 \end{bmatrix} = \begin{bmatrix} 10 & 4 \end{bmatrix}$$




Moscow

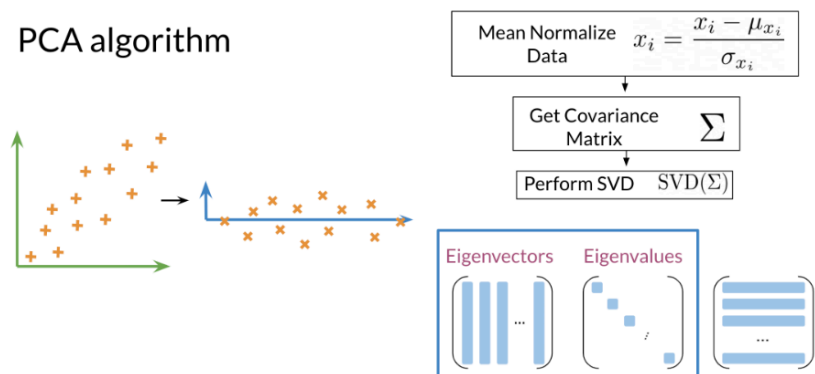
### PCA:

- Dimensionality Reduction
- Unsupervised

### Principal Component Analysis



## PCA algorithm



### Steps to Compute PCA:

- Mean normalize your data
- Compute the covariance matrix
- Compute SVD on your covariance matrix. This returns  $[USV] = \text{svd}(\Sigma)$ . The three matrices U, S, V are drawn above. U is labelled with eigenvectors, and S is labelled with eigenvalues.
- You can then use the first n columns of vector  $U$ , to get your new data by multiplying  $XU[:, 0 : n]$ .

## Machine Translation:

### Transforming word vectors:

$$\begin{pmatrix} \text{[ "cat" vector ]} \\ \text{[ ... vector ]} \\ \text{[ "zebra" vector ]} \end{pmatrix} \mathbf{X} \mathbf{R} \approx \mathbf{Y} \begin{pmatrix} \text{[ "chat" vecteur ]} \\ \text{[ ... vecteur ]} \\ \text{[ "zébrasse" vecteur ]} \end{pmatrix} \mathbf{Y}$$

subsets of the full vocabulary

### Steps required to learn $R$ :

- Initialize  $R$
- For loop

$$Loss = \|XR - Y\|_F$$

$$g = \frac{d}{dR} Loss$$

$$R = R - \alpha * g$$

### Frobenius Norm:

$$\|\mathbf{XR} - \mathbf{Y}\|_F$$

$$\mathbf{A} = \begin{pmatrix} 2 & 2 \\ 2 & 2 \end{pmatrix}$$

$$\|\mathbf{A}_F\| = \sqrt{2^2 + 2^2 + 2^2 + 2^2}$$

$$\|\mathbf{A}_F\| = 4$$

$$\|\mathbf{A}\|_F \equiv \sqrt{\sum_{i=1}^m \sum_{j=1}^n |a_{ij}|^2}$$

**Note:** For Simplicity, we can minimize the square of frobenius norm.

$$Loss = \|\mathbf{XR} - \mathbf{Y}\|_F^2$$

$$g = \frac{d}{dR} Loss = \frac{2}{m} (\mathbf{X}^T (\mathbf{XR} - \mathbf{Y}))$$

Here, m is the number of words or rows in the R matrix.

**K-nearest neighbours:**

**Hash Tables:**


0	1	2	3	4	5	6	7	8	9
100				14			17		
10							97		

Hash function (vector)  $\longrightarrow$  Hash value


Hash value = vector % number of buckets

```
def basic_hash_table(value_l, n_buckets):
    def hash_function(value_l, n_buckets):
        return int(value_l) % n_buckets
    hash_table = {i: [] for i in range(n_buckets)}
    for value in value_l:
        hash_value = hash_function(value, n_buckets)
        hash_table[hash_value].append(value)
    return hash_table
```

**Locality sensitive hashing:**



**P is the perpendicular to the plane**



$$P_1 v^T = 3, \text{sign}_1 = +1, h_1 = 1$$

$$P_2 v^T = 5, \text{sign}_2 = +1, h_2 = 1$$

$$P_3 v^T = -2, \text{sign}_3 = -1, h_3 = 0$$

$$\text{hash} = 2^0 \times h_1 + 2^1 \times h_2 + 2^2 \times h_3$$

$$= 1 \times 1 + 2 \times 1 + 4 \times 0$$

$$= 3$$

```
def hash_multiple_plane(P_l,v):
    hash_value = 0
    for i, P in enumerate(P_l):
        sign = side_of_plane(P,v)
        hash_i = 1 if sign >=0 else 0
        hash_value += 2**i * hash_i
    return hash_value
```

**Approximate nearest neighbors:**

```
num_dimensions = 2 #300 in assignment
num_planes = 3 #10 in assignment

random_planes_matrix = np.random.normal(
    size=(num_planes,
          num_dimensions))
```

```
array([[ 1.76405235  0.40015721]
       [ 0.97873798  2.2408932 ]
       [ 1.86755799 -0.97727788]])
```

```
v = np.array([[2,2]])
```

```
def side_of_plane_matrix(P,v):
    dotproduct = np.dot(P,v.T)
    sign_of_dot_product = np.sign(dotproduct)
    return sign_of_dot_product
```

```
num_planes_matrix = side_of_plane_matrix(
    random_planes_matrix,v)
```

```
array([[1.]
       [1.]
       [1.]])
```

See notebook for calculating the hash value!

**Searching Document:**

```
word_embedding = {"I": np.array([1,0,1]),
                  "love": np.array([-1,0,1]),
                  "learning": np.array([1,0,1])}

words_in_document = ['I', 'love', 'learning']
document_embedding = np.array([0,0,0])

for word in words_in_document:
    document_embedding += word_embedding.get(word,0)

print(document_embedding)
array([1 0 3])
```